

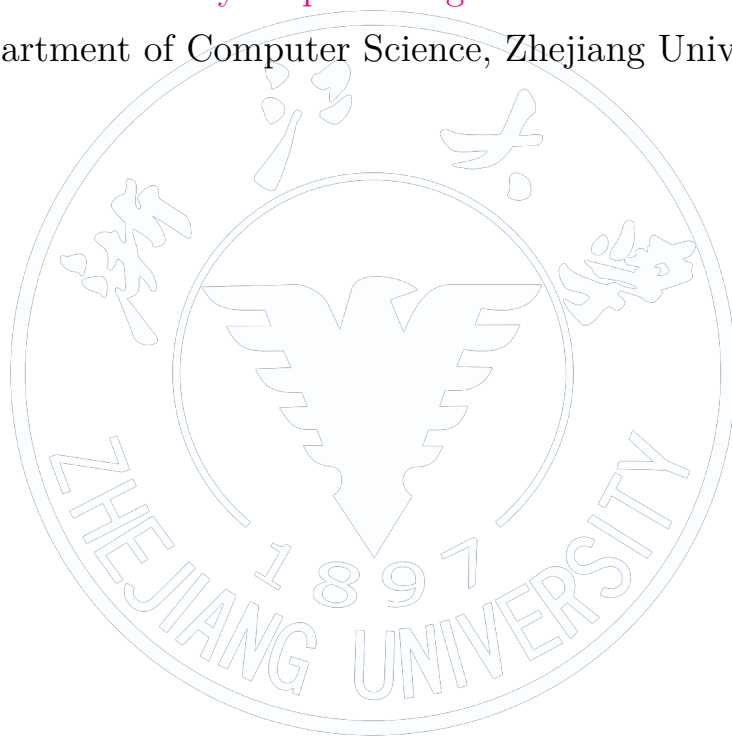
計算機組織

Computer Architecture

TZU-CHUN HSU¹

¹vm3y3rmp40719@gmail.com

¹Department of Computer Science, Zhejiang University



2022 年 3 月 12 日
Version 5.0.0

Disclaimer

本文「計算機組織」為台灣研究所考試入學的「計算機組織」考科使用，內容主要參考張凡先生的二本計算機組織參考書 [1][2]，以及 wjungle 網友在 PTT 論壇上提供的資料結構筆記 [3]。

本文作者為 TZU-CHUN HSU，本文及其 L^AT_EX 相關程式碼採用 MIT 協議，更多內容請訪問作者之 GITHUB 分頁 [Oscarshu0719](#)。

MIT License

Copyright (c) 2020 TZU-CHUN HSU

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

1 Overview

1. 本文頁碼標記依照實體書 [1][2] 的頁碼。
2. TKB 筆記 [3] 章節頁碼：

Chapter	Page No.
1	1
2	27
3	81
4	101
5	119
6	165
7	221
8	238

3. 省略第一章重點十四，第二章重點五、六、十，第六章重點十四，第七章只看重點一到四及十，第八章重點五一致性協定範例。

2 Summary

1. Theorem (10) Endianness:

- Big Endian: 最左邊或 MSB 放在最低 address, e.g. MIPS。
- Little Endian: 最右邊或 LSB 放在最低 address, e.g. x86。

2. Theorem (28, 47, 59)

- `srl/sll rd, rt, shamt # rs = 5'0`
- `lw/sw rt, imm(rs)`
- `beq/bne rs, rt, addr`
- `addi rt, rs, imm`
- `lb`: Load 最高 address (LSB) 到最高 address (LSB), `sb`: Store 最高 address (LSB) 到最低 address (MSB)。
- `jr rs # rt = rd = shamt = 5'0`: R-type

3. Theorem (62)

```
int fact (int n) {  
    if (n < 1)  
        return 1;  
    else  
        return (n * fact(n - 1));  
}
```

fact:

```
addi $sp, $sp, -8  
sw $ra, 4($sp)  
sw $a0, 0($sp)  
slti $t0, $a0, 1  
beq $t0, $zero, L1  
addi $v0, $zero, 1  
addi $sp, $sp, 8  
jr $ra
```

L1:

```

addi $a0, $a0, -1
jal fact
lw $a0, 0($sp)
lw $ra, 4($sp)
addi $sp, $sp, 8
mul $v0, $a0, $v0
jr $ra

```

4. Theorem (190) 浮點數:

Single precision		Double precision		Representation
Exponent	Fraction	Exponent	Fraction	
0	0	0	0	± 0
0	$\neq 0$	0	$\neq 0$	\pm denormalized number
1 ~ 254	\times	1 ~ 2046	\times	\pm floating-point number
255	0	2047	0	$\pm \infty$
255	$\neq 0$	2047	$\neq 0$	NaN

5. Theorem (214, 215) Overflow detection:

- 有號數:

```

addu $t0, $t1, $t2
xor $t3, $t1, $t2
slt $t3, $t3, $zero # $t3 = 1 if sign differs
bne $t3, $zero, NO_OVERFLOW
xor $t3, $t0, $t1 # Check if the sum sign differs
slt $t3, $t3, $zero
bne $t3, $zero, OVERFLOW

```

- 無號數:

```

addu $t0, $t1, $t2
nor $t3, $t1, $zero #  $2^{32} - \$t1 - 1$ 
sltu $t3, $t3, $t2 #  $2^{32} - \$t1 - 1 < \$t2 \rightarrow 2^{32} - 1 < \$t1 + \$t2$ 
bne $t3, $zero, OVERFLOW

```

6. Theorem (371) 只有 `jump` 和 `MemoReg` 上面 1 下面 0, 其他皆相反。

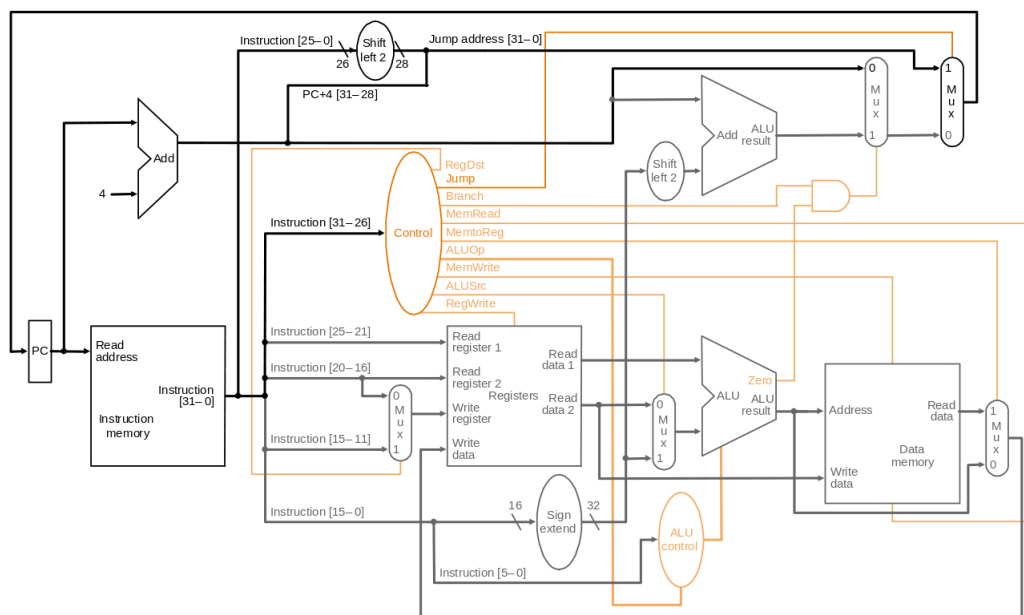


Figure 1: Single-cycle CPU with jump and branch.

7. Theorem (384)

Instruction	ALUOp1	ALUOp2
lw/sw	0	0
beq	×	1
R-type	1	×

8. Theorem (441) 原始 pipeline 設計:

- `beq` 在 *MEM* 決定是否要跳。
- `RegDst` 在 *EX*。




Figure 2: Original pipeline.

9. Theorem (450, 455, 457, 458) Data hazards:

- Forwarding: Combinational units, 放在 EX 因為 ALU。

if $(EX/MEM.RegWrite \wedge (EX/MEM.Rd \neq 0) \wedge (EX/MEM.Rd = ID/EX.Rs/Rt))$
 ForwardA/B = 10

Listing 1: EX hazard.

if $(MEM/WB.RegWrite \wedge (MEM/WB.Rd \neq 0) \wedge (\neg EX_hazard) \wedge (MEM/WB.Rd = ID/EX.Rs/Rt))$
 ForwardA/B = 01

Listing 2: MEM hazard.

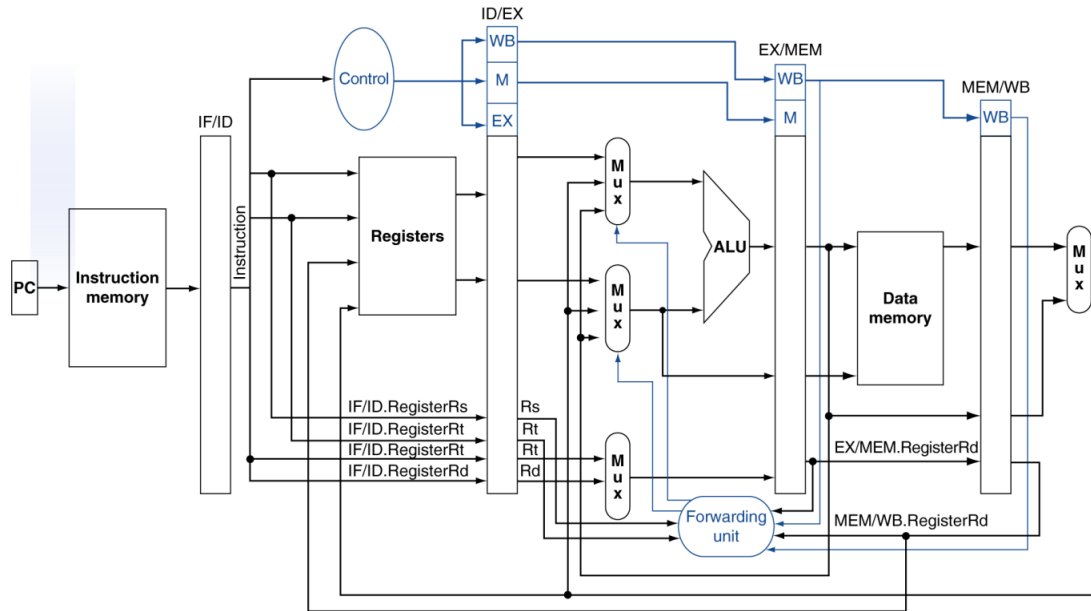


Figure 3: Pipeline with forwarding.

- Stall:

```

if (ID/EX.MemRead  $\wedge$  (ID/EX.Rt = IF/ID.Rs/Rt))
  IF/ID.Write := 0
  PC.Write := 0

```

Listing 3: Stall.

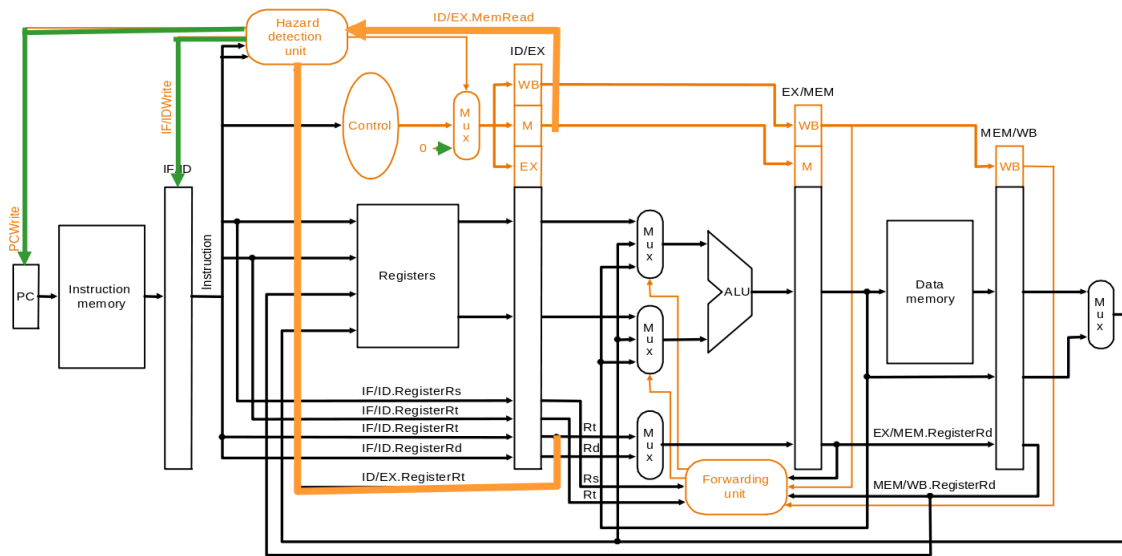


Figure 4: Pipeline with hazard detection and forwarding units.

10. Theorem (478, 487, 494, 559) Control hazards:

- 若分支指令與前一個 ALU 指令或前面第二個 lw 有 data dependency，必須 stall 1 CC。
- 分支指令通過 xor 再 nor 比較是否相同。




Figure 5: Pipeline with hazard detection, forwarding units and flush.

- Delayed branch:
 - NOT suitable for deep pipeline.
 - From before: 最佳方法，不管跳或不跳皆提升。
 - From target: 用於 branch 發生機率高，有跳才提升。
 - From fall through: 用於 branch 發生機率低，不跳才提升。

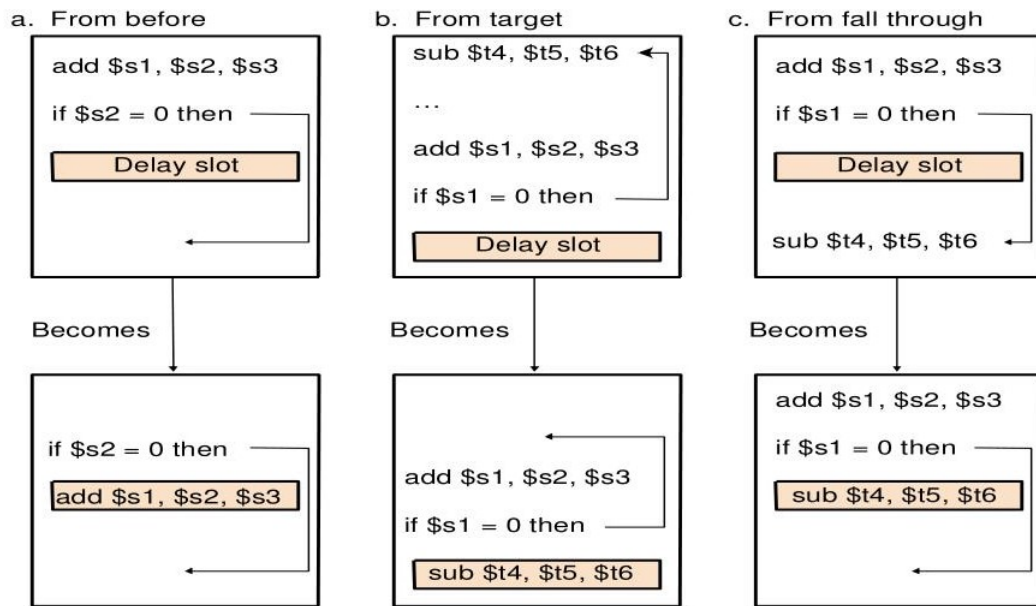


Figure 6: Example of delayed branch.

11. Theorem (499, 505, 511)

- Intel IA-64 (EPIC):
 - 支援利用 compiler 開發的平行度。
 - 可以猜測，並利用 if-else 取代 branch。
 - Registers 比 MIPS 多很多。
 - Instruction group is a sequence of instructions which does **NOT** have data dependency and can be executed parallelly.
- Speculation 錯誤復原:
 - 軟體提供修補程式。
 - 硬體 CPU 將猜測結果暫時儲存，若正確，則將猜測結果寫回 register 或 memory，否則 flush buffer。

Advanced pipeline		
Technique	Hardware	Software
Branch prediction	✓	✓
Speculation	✓	✓
Intel IA-64 (EPIC)	✓	✓
Register renaming	✓	✓
Prediction		✓

12. **Theorem (515)** MIPS exception handling:

- Flush the instruction and let all preceding instructions complete if they can.
- 利用 cause register 儲存 exception 原因。
- 將造成 exception 的 instruction memory address 存在 EPC ($PC + 4$)。
- 使用 entry point switch to kernel。
- Exception handling routine 須將 $PC - 4$ 。

13. **Theorem (195)** Non-blocking cache:

- Does **NOT** allow **miss under hit** to hide miss latency.
- **Miss under miss** allows multiple outstanding cache misses.
- Allow a **load** instruction to access the cache if the previous **load** is a cache miss.

14. **Theorem (213, 278, 289)** RAID:

- RAID 2: Hamming code, Write 需要讀取所有 disks, 從新計算 Hamming code 並寫入 ECC disks, 效率差, $2n - 1$ disks。
- RAID 3:
 - Reliability 和 RAID 2 相同。
 - 不做備份, 花費較多時間恢復 data, $n + 1$ disks。
 - 當 1 個 disk 出錯可救回來, 多個則否。
 - Availability cost 為 $\frac{1}{N}$, 其中 N 為 protection group disks 數量。
 - Parity 集中存放一個 disk。
- RAID 4:
 - 只對 protection group 其中一 disk 做 small reads。
 - $n + 1$ disks, parity 集中存放一個 disk。

- 當 1 個 disk 出錯可救回來，多個則否。
- RAID 5:
 - Write 就不會有單一 disk 瓶頸。
 - $n + 1$ disks, parity 被分散到所有 disks。
 - 可允許 1 個 disk 故障。
- RAID 6:
 - 與 RAID 5 相比，增加第二個獨立的 parity block。
 - 通常通過硬體實現。
 - $n + 2$ disks。
 - 可允許 2 個 disk 故障。
- RAID 3 has **worst** throughput for **small writes**.
- RAID 3 has **best small writes latency**.
- RAID 3, 4, 5 have same throughput for **large writes**.
- RAID 1 can **NOT** have **small writes** in parallel.
- RAID 3 can **NOT** have **small writes or reads** in parallel.
- RAID 4, 5 perform same for parallel **small reads and writes**.
- RAID 4 does **NOT** have better **big reads** performance than RAID 3.
- RAID 1+0 has better **write throughput** than RAID 0+1.

15. Theorem (320)

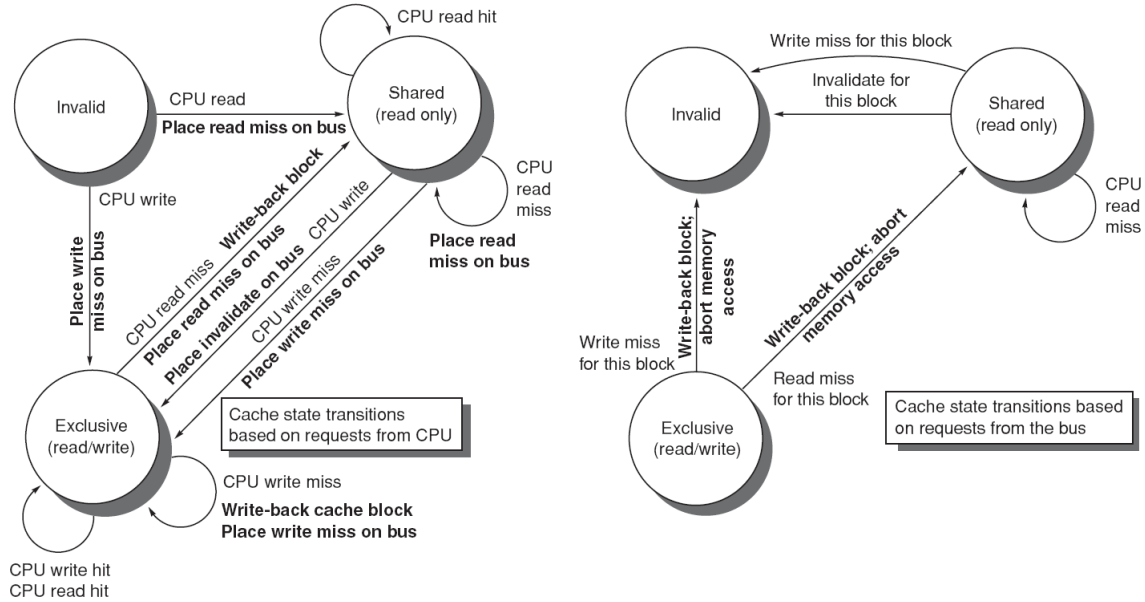


Figure 7: Snooping states.

16. **Theorem (325) Multithreading:**

- Coarse-grained multithreading: Switch threads only on costly stalls, such as L2 cache misses. Pipeline **start-up** costs.
- Fine-grained multithreading: Switch between threads on each instruction packs. It can hide the throughput losses.
- SMT: **ILP** and **TLP**; coarse-grained and fine-grained: **TLP only**.

17. **Theorem (334) Network topology:**

- Performance measure:
 - Network bandwidth.
 - Bisection bandwidth: 平均切為二所減少的 bandwidth, 越高容錯力越高。
 - Diameter: 任兩點最短路的最大值, 越低越好。
 - Nodal degree: CPU degree, 越高容錯力越高。
- Omega network hardware: $2n \log_2 n$.
- Crossbar network hardware: n^2 .

18. **Theorem () NAS vs SAN:**

- NAS operates at **file** level while SAN operates at **block** level.

- CIFS/SMB and NFS are examples of NAS.
- SAN is often the preferred choice over NAS.
- Almost any machine running Microsoft Windows with LAN connectivity can be configured to access a NAS.

19. **Theorem ()** Log-structured file system:

- 將要 write 的 data 合成一串，再一次 write。
- Read 都在 cache，因為 cache 夠大。
- Disk access 的 seek 和 rotation 是 bottleneck，sequential access 比 random access 好。

20. **Theorem ()**

- Meltdown: read arbitrary kernel memory, and it does **NOT** rely on software vulnerabilities.
- Spectre: Making other applications to access arbitrary contents in memory.
- Both belongs to **side channel attacks**.
- Does **NOT** leave records in traditional log.
- Hard for antivirus software to detect them.
- Processors which are able to implement out-of-order execution is risky.
- IA-64 is immune to Spectre and Meltdown.

21. **Theorem ()** Power:

- CMOS does **NOT** consume power when it's **static** ($power_{static} = 0$), so it can decrease **frequency** to save power.
- **Static** power dissipation occurs because of leakage current that flows even when a transistor is **off**.
- Computers at **lower utilization** does **NOT** use less power proportionally.
- The main reason for the switch from high-performance uniprocessors to multiprocessors with simpler cores and lower clock rates in recent years is the **power limit** and **memory gap**.

22. **Theorem ()** Cache:

- L1 data cache is usually separated from L1 instruction cache to **increase bandwidth**.
- Data cache is usually deployed at **MEM** stage.
- In modern processors, **L1** data and instruction caches are split, but L2 does **NOT**. Both L1 and L2 caches are **write-back**.

23. **Theorem ()** Branch prediction:

- Branch target buffer is used by **CPU**, which is checked at **IF** stage.
- Branch prediction buffer is good to predict the **branch outcome**, but it does **NOT** help in predicting the **branch target**.
- Indirect branch prediction: Dynamic: hybrid predictor; Static: Neural branch predictor.
- **Virtual program counter prediction** is often used to predict **conditional/unconditional indirect** branch, which treats indirect branches as **multiple conditional branches**.

24. **Theorem ()** Hazards:

- **Memory hazard** do NOT cause stall, e.g. **sw** after **lw**.
- Control hazards can **NOT** be avoided.

25. **Theorem ()** Page table:

- In hash-based page tables using **linked list** to solve collision, **each element** contains a frame number and a page number.
- MIPS uses **two** separated page tables and two limit registers, one for **stack** and the other for **heap**.

26. **Theorem ()** Arithmetic:

- Increasing number of **used sticky bits** do NOT improve accuracy.
- Conversion from single-precision to double-precision causes loss of precision.
- Ripple Carry Adder: Critical path delay is $2N$ gate delay (carry out), and sum delay is $2N + 1$ gate delay (actual sum).
- (109NYCU109-15d) Converting an integer variable to a **single** precision FP number will lose precision, but **double** precision does **NOT**.

27. **Theorem ()** Multi-threading:

- GPGPU usually runs **SPMT** (Single Program Multiple Thread), and GPU runs **SIMT**.
- Vector processors need **less bandwidth** than conventional processors.
- GPUs do **NOT** rely on **multilevel caches**.

28. Theorem ()

- Out-of-order execution in **cache** level do NOT fail.
- Program is a **passive** entity, process is an **active** entity.
- Multiple-cycles CPU requires **minimum function units**.
- Compiler identifies **basic blocks** for code optimization.
- To form the machine code, the value of label of branch instructions is computed by **linker** when the label is an **external** reference.
- **NOT** each computer support **direct addressing mode**.
- **Conflict** misses do **NOT** occur in **fully associative** caches.
- **MIPS** and **ARM** use **memory-mapped I/O**.
- Writes are much **slower** than reads for flash. NAND flash is **cheaper** than NOR flash.
- (109NYCU-20a) (**FALSE**) It's NOT necessary to speculate two instructions with a **store preceding a load** that reference the same memory location, where this kind of speculation allows to perform a load before a store, since sometimes we can NOT determine if two locations referenced by a store and a load are the same due to the pointer aliasing problem.
- Difficulty to handle **exceptions** (from most difficult to simplest):

Superscalar
Speculative
Out-of-order
Pipelined
Single-issue in-order processor
Hierarchical data caches

- Difficulty to handle **interrupts** (from most difficult to simplest):

GPGPU
Containers
Virtual machines
Hyper-threaded processor
Superscalar
Pipelined



References

- [1] 張凡. 計算機組織與結構重點直擊（上）. 鼎茂圖書出版股份有限公司, 3 edition, 2019.
- [2] 張凡. 計算機組織與結構重點直擊（下）. 鼎茂圖書出版股份有限公司, 3 edition, 2019.
- [3] wjungle@ptt. 計算機組織 @tkb 筆記. <https://drive.google.com/file/d/0B8-2o6L73Q2VUkpEMWVLb1pRZE0/view?usp=sharing>, 2017.

